

Security Pillar

AWS Well-Architected Framework

July 2018



Notices

This document is provided for informational purposes only. It represents AWS's current product offerings and practices as of the date of issue of this document, which are subject to change without notice. Customers are responsible for making their own independent assessment of the information in this document and any use of AWS's products or services, each of which is provided "as is" without warranty of any kind, whether express or implied. This document does not create any warranties, representations, contractual commitments, conditions or assurances from AWS, its affiliates, suppliers or licensors. The responsibilities and liabilities of AWS to its customers are controlled by AWS agreements, and this document is not part of, nor does it modify, any agreement between AWS and its customers.

Contents

Introduction	1
Security	1
Design Principles	2
Definition	3
Identity and Access Management	3
Protecting AWS Credentials	3
Fine-Grained Authorization	6
Detective Controls	7
Capture and Analyze Logs	7
Integrate Auditing Controls with Notification and Workflow	10
Infrastructure Protection	12
Protecting Network and Host-Level Boundaries	13
System Security Configuration and Maintenance	15
Enforcing Service-Level Protection	16
Data Protection	17
Data Classification	18
Encryption/Tokenization	19
Protecting Data at Rest	21
Protecting Data in Transit	23
Data Backup/Replication/Recovery	24
Incident Response	26
Clean Room	26
Conclusion	28
Contributors	28
Further Reading	29
Document Revisions	29

Abstract

The focus of this paper is the security pillar of the [Well-Architected Framework](#). It provides guidance to help you apply best practices in the design, delivery, and maintenance of secure AWS environments.

Introduction

The [AWS Well-Architected Framework](#) helps you understand the pros and cons of decisions you make while building systems on AWS. By using the Framework you will learn architectural best practices for designing and operating reliable, secure, efficient, and cost-effective systems in the cloud. It provides a way for you to consistently measure your architectures against best practices and identify areas for improvement. We believe that having well-architected systems greatly increases the likelihood of business success.

The framework is based on five pillars:

- Operational Excellence
- Security
- Reliability
- Performance Efficiency
- Cost Optimization

This paper focuses on the security pillar and how to apply it to your solutions. Ensuring security can be challenging in traditional on-premises solutions due to the use of manual processes, eggshell security models, and insufficient auditing. By adopting the practices in this paper you can build architectures that protect data and systems, control access, and respond automatically to security events.

This paper is intended for those in technology roles, such as chief technology officers (CTOs), architects, developers, and operations team members. After reading this paper, you will understand AWS best practices and strategies to use when designing cloud architectures for security. This paper doesn't provide implementation details or architectural patterns; however, it does include references to appropriate resources for this information.

Security

The security pillar encompasses the ability to protect information, systems, and assets while delivering business value through risk assessments and mitigation strategies. This paper will provide in-depth, best-practice guidance for architecting secure systems on AWS.

Design Principles

In the cloud, there are a number of principles that can help you strengthen your system security:

- **Implement a strong identity foundation:** Implement the principle of least privilege and enforce separation of duties with appropriate authorization for each interaction with your AWS resources. Centralize privilege management and reduce or even eliminate reliance on long-term credentials.
- **Enable traceability:** Monitor, alert, and audit actions and changes to your environment in real time. Integrate logs and metrics with systems to automatically respond and take action.
- **Apply security at all layers:** Rather than just focusing on protection of a single outer layer, apply a defense-in-depth approach with other security controls. Apply to all layers (e.g., edge network, VPC, subnet, load balancer, every instance, operating system, and application).
- **Automate security best practices:** Automated software-based security mechanisms improve your ability to securely scale more rapidly and cost effectively. Create secure architectures, including the implementation of controls that are defined and managed as code in version-controlled templates.
- **Protect data in transit and at rest:** Classify your data into sensitivity levels and use mechanisms, such as encryption, tokenization, and access control where appropriate.
- **Keep people away from data:** Create mechanisms and tools to reduce or eliminate the need for direct access or manual processing of data. This reduces the risk of loss or modification and human error when handling sensitive data.
- **Prepare for security events:** Prepare for an incident by having an incident management process that aligns to your organizational requirements. Run incident response simulations and use tools with automation to increase your speed for detection, investigation, and recovery.

Definition

Security in the cloud is composed of five areas:

1. Identity and access management
2. Detective controls
3. Infrastructure protection
4. Data protection
5. Incident response

The AWS Shared Responsibility Model enables organizations that adopt the cloud to achieve their security and compliance goals. Because AWS physically secures the infrastructure that supports our cloud services, as an AWS customer you can focus on using services to accomplish your goals. The AWS Cloud also provides greater access to security data and an automated approach to responding to security events.

Identity and Access Management

Identity and access management are key parts of an information security program, ensuring that only authorized and authenticated users are able to access your resources, and only in a manner that you intend. For example, you should define principals (that is, users, groups, services, and roles that take action in your account), build out policies aligned with these principals, and implement strong credential management. These privilege-management elements form the core of authentication and authorization.

In AWS, there are a number of different approaches to consider when addressing identity and access management. The following sections describe how to use these approaches:

- Protecting AWS credentials
- Fine-grained authorization

Protecting AWS Credentials

The careful management of access credentials is the foundation of how you will secure your resources in the cloud. Every interaction you make with AWS will

be authenticated, so establishing appropriate credential management practices and patterns allows you to tie the use of AWS to your workforce lifecycle and ensure that only the appropriate parties take action in your account.

When you open an AWS account, the identity you begin with has access to all AWS services and resources in that account. You use this identity to establish less-privileged users and role-based access in the AWS Identity and Access Management (IAM) service. However, this initial account (known as the root user) isn't intended for everyday tasks, and these credentials should be carefully protected using multi-factor authentication (MFA) and by deleting any access keys upon completion of the initial account setup.

For the root user, you should follow the best practice of only using this login to create another, initial set of IAM users and groups for longer-term identity management operations. These privileged IAM users – carefully monitored and constrained – can be used to assume roles in one or many accounts you own. You may choose to establish trust with existing identity providers using federation (via SAML 2.0 or web identities) already tied to your organization's workforce source of record. Using federation reduces the need to create users in IAM while leveraging the existing identities, credentials, and role-based access you might already have established in your organization.

For all IAM users, you should apply appropriate policies enforcing the use of strong authentication. You should set a password policy on the AWS account that requires a minimum length and complexity for passwords associated with IAM users. You should also set a mandatory rotation policy requiring IAM users to change their passwords at regular intervals. For all IAM users with passwords permitting access to the AWS Management Console, you should also require the use of MFA.

IAM users might also require access to the AWS APIs directly from command-line tools (CLI) or by using software development kits (SDKs). In these cases, where federation might not be practical, an access key ID and secret access key can be issued and used in place of, or in addition to, a password. IAM roles should be used to grant permission, in this case with permissions granted on the role. The IAM user can only assume that role with enforcement of MFA. These credentials should be carefully protected and exchanged for temporary credentials whenever possible. Take extra care to avoid storing access and secret

keys in improperly secured locations or inadvertently committing them to source code repositories.

For use cases where federation or IAM roles might not be practical, such as different service-to-service authentication scenarios, you can use IAM instance profiles for Amazon Elastic Compute Cloud (Amazon EC2) instances and the AWS Security Token Service (AWS STS) to generate and manage temporary credentials used in software that must authenticate to AWS APIs.

Key AWS Services

The key AWS service that supports the protection of credentials is IAM. This service allows you to manage credentials and the policies applied to them. The following services and features are also important:

- **AWS STS** lets you request temporary, limited-privilege credentials for authentication with other AWS APIs.
- **IAM instance profiles for EC2 instances** allow you to leverage the Amazon EC2 metadata service and managed, temporary credentials for accessing other AWS APIs.

Resources

Refer to the following resources to learn more about AWS best practices for protecting your AWS credentials.

Videos

- [IAM Best Practices to Live By](#)

Documentation

- [The AWS Account Root User](#)
- [AWS Account Root User Credentials vs. IAM User Credentials](#)
- [IAM Best Practices](#)
- [Setting an Account Password Policy for IAM Users](#)
- [Managing Access Keys for IAM Users](#)
- [Using Instance Profiles](#)

- [Temporary Security Credentials](#)

Fine-Grained Authorization

Establishing a principle of least privilege ensures that authenticated identities are only permitted to perform the most minimal set of functions necessary to fulfill a specific task, while balancing usability and efficiency. Operating on this principle limits the blast radius - or potential impact - of inappropriate use of valid credentials. The principle of least privilege allows you to enforce separation of duties for oversight and governance, and makes auditing the entitlements to your resources much simpler.

Your organization should define roles and responsibilities for users and applications interacting with AWS services and implement fine-grained authorization for enforcement of these roles.

Fine-grained authorization is implemented in AWS using IAM roles and policies. A role is an IAM “principal” assumed by a user or another AWS service and is assigned temporary credentials scoped to a limited set of permissions. IAM policies are documents that formally state one or more permissions. Policies are attached to users, groups, and roles to create a very robust access management framework.

AWS Organizations should be used to centrally manage AWS accounts. This service allows grouping accounts into organizational units (OUs). Service control policies (SCPs) can be used to centrally control AWS services across multiple AWS accounts.

Key AWS Services

The key AWS service supporting fine-grained authorization is IAM, which offers a highly flexible policy language used to allow or deny actions, set conditions on these actions, and constrain actions to specific principals or resources. The following service is also important:

- **AWS Organizations** lets you centrally manage and enforce policies for multiple AWS accounts.

Resources

Refer to the following resources to learn more about AWS best practices for fine-grained authorization.

Videos

- [IAM Policy Ninja](#)

Documentation

- [Working with Policies](#)
- [Delegating Permissions to Administer IAM Users, Groups, and Credentials](#)
- [Remove unnecessary credentials](#)
- [Assuming a role in the CLI with MFA](#)
- [AWS Organizations](#)

Detective Controls

You can use detective controls to identify a potential security threat or incident. They are an essential part of governance frameworks and can be used to support a quality process, a legal or compliance obligation, and for threat identification and response efforts. There are different types of detective controls. For example, conducting an inventory of assets and their detailed attributes promotes more effective decision making (and lifecycle controls) to help establish operational baselines. You can also use internal auditing, an examination of controls related to information systems, to ensure that practices meet policies and requirements and that you have set the correct automated alerting notifications based on defined conditions. These controls are important reactive factors that can help your organization identify and understand the scope of anomalous activity.

In AWS, there are a number of approaches to consider when addressing detective controls. The following sections describe how to use these approaches:

- Capture and analyze logs
- Integrate auditing controls with notification and workflow

Capture and Analyze Logs

In traditional data center architectures, aggregating and analyzing logs typically requires installing agents on servers, carefully configuring network appliances

to direct log messages at collection points, and forwarding application logs to search and rules engines. Aggregation in the cloud is much easier due to two capabilities.

First, asset management is easier because assets and instances can be described programmatically without depending on agent health. For example, instead of manually updating an asset database and reconciling it with the real install base, you can reliably gather asset metadata with just a few API calls. This data is far more accurate and timely than using discovery scans, manual entries into a configuration management database (CMDB), or relying on agents that might stop reporting on their state.

Second, you can use native, API-driven services to collect, filter, and analyze logs instead of maintaining and scaling the logging backend yourself. Pointing your logs at a bucket in an object store, or directing events to a real-time log processing service, means that you can spend less time on capacity planning and availability of the logging and search architecture.

In AWS, a best practice is to customize the delivery of AWS CloudTrail and other service-specific logging to capture API activity globally and centralize the data for storage and analysis. You can direct CloudTrail logs to Amazon CloudWatch Logs or other endpoints so that you can obtain events in a consistent format across compute, storage, and applications.

For instance-based and application-based logging that doesn't originate from AWS services, you can still use agent-based tools to collect and route events. You can use Amazon CloudWatch Logs to monitor, store, and access your log files from Amazon Elastic Compute Cloud (Amazon EC2) instances, AWS CloudTrail, Amazon Route 53, and other sources. Using services and features such as AWS CloudFormation, AWS Systems Manager, or Amazon EC2 user data, system administrators can ensure that instances always have agents installed.

Equally important to collecting and aggregating logs is the ability to extract meaningful insight from the great volumes of log and event data generated by modern, complex architectures. See the *Monitoring* section of [The Reliability Pillar](#) whitepaper for more detail. Architects should consider detective controls end-to-end. You should not simply generate and store logs. Your information security function needs robust analytics and retrieval capabilities to provide

insight into security related activity. For intelligent threat detection, you can use a service such as Amazon GuardDuty to continuously monitor events from AWS CloudTrail, Amazon VPC Flow Logs, and DNS logs.

Key AWS Services

The key AWS service that supports capturing key activities is CloudTrail, which provides rich detail about API calls made in your AWS account. The following services and features are also important:

- **Amazon GuardDuty** is a managed threat detection service that continuously monitors for malicious or unauthorized behavior to help you protect your AWS accounts and workloads.
- **AWS Config** provides you with an AWS resource inventory, configuration history, and configuration change notifications to enable security and governance.
- **Amazon CloudWatch Logs** allows you to centralize logs into streams, natively integrating with features and services like Amazon VPC Flow Logs and CloudTrail. CloudWatch Logs scales to ingest logs without the need to manage infrastructure.
- **Amazon Simple Storage Service (S3) and Amazon Glacier** can be used to centralize storage and long-term archiving of log data.
- **Amazon Athena** can be used to analyze logs, such as CloudTrail logs, to help you identify trends and further isolate activity by attribute, such as source IP address or user.

Resources

Refer to the following resources to learn more about AWS best practices for capturing and analyzing logs.

Documentation

- [Setting up Amazon GuardDuty](#)
- [Getting started: Amazon CloudWatch Logs](#)
- [Configuring Athena to analyze CloudTrail logs](#)

Integrate Auditing Controls with Notification and Workflow

Security operations teams rely on the collection of logs and the use of search tools to discover potential events of interest, which may indicate unauthorized activity or unintentional change. However, simply analyzing collected data and manually processing information is insufficient to keep up with the volume of information flowing from modern, complex architectures. Analysis and reporting alone don't facilitate the assignment of the right resources to work an event in a timely fashion. A best practice for building a mature security operations team is to deeply integrate the flow of security events and findings into a notification and workflow system such as a ticketing system, a bug/issue system, or other security information and event management (SIEM) system. This takes the workflow out of email and static reports, allowing you to route, escalate, and manage events or findings. Many organizations are now integrating security alerts into their chat/collaboration and developer productivity platforms.

This best practice applies not only to security events generated from log messages depicting user activity or network events, but from changes detected in the infrastructure itself. The ability to detect change, determine whether a change was appropriate, and then route this information to the correct remediation workflow is essential in maintaining and validating a secure architecture.

In AWS, routing events of interest and information reflecting potentially unwanted changes into a proper workflow is done using Amazon CloudWatch Events. This service provides a scalable rules engine designed to broker both native AWS event formats (such as CloudTrail events), as well as custom events you can generate yourself. You build rules that parse events, transform them if necessary, and then route such events to targets such as an AWS Lambda function, Amazon Simple Notification Service (Amazon SNS) notification, or other targets.

Detecting change and routing this information to the correct workflow can be accomplished using AWS Config rules. AWS Config detects changes to in-scope services and generates events that can be parsed using AWS Config rules for rollback, enforcement of compliance policy, and forwarding of information to

systems, such as change management platforms and operational ticketing systems.

Reducing the number of security misconfigurations introduced into a production environment is critical, so the more quality control and reduction of defects you can perform in the build process, the better. Modern continuous integration and continuous deployment (CI/CD) pipelines should be designed to test for security issues whenever possible. Using Amazon Inspector, you can perform configuration assessments for known common vulnerabilities and exposures (CVEs), assess your instances against security benchmarks, and fully automate the notification of defects. Amazon Inspector runs on production instances or in a build pipeline, and it notifies developers and engineers when findings are present. You can access findings programmatically and direct your team to backlogs and bug-tracking systems.

Key AWS Services

The key AWS service that supports integrating auditing controls into notification and workflow systems is CloudWatch Events, which allows you to route events to a powerful rules engine. Rules then examine incoming events, parse the incoming values, and properly route the event to any number of targets, such as email or mobile devices, ticketing queues, and issue management systems. The following services and features are also important:

- **AWS Config Rules** enables you to create rules that automatically check the configuration of AWS resources recorded by AWS Config.
- **Amazon CloudWatch** and **CloudWatch Logs** must be enabled to facilitate collection of events and routing with CloudWatch Events.
- **Amazon CloudWatch API and AWS SDKs** can be used to create custom events in your own applications and inject them into CloudWatch Events for rule-based processing and routing.
- **Amazon Inspector** offers a programmatic way to find security defects or misconfigurations in your operating systems and applications. Because you can use API calls to access both the processing of assessments and the results of your assessments, integration of the findings into workflow and notification systems is simple. DevOps teams can integrate Amazon Inspector into their CI/CD pipelines and use it to identify any pre-existing issues or when new issues are introduced.

Resources

Refer to the following resources to learn more about AWS best practices for integrating auditing controls with notification and workflow.

Videos

- [Best Practices for Managing Security Operations on AWS](#)
- [Achieving Continuous Compliance using AWS Config](#)

Documentation

- [Amazon CloudWatch Events](#)
- [AWS Config Rules](#)
- [AWS Config Rules Repository \(open source on GitHub\)](#)
- [Amazon Inspector](#)

Infrastructure Protection

Infrastructure protection encompasses control methodologies, such as defense in depth, necessary to meet best practices and organizational or regulatory obligations. Use of these methodologies is critical for successful, ongoing operations in either the cloud or on premises.

Infrastructure protection is a key part of an information security program. It ensures that systems and services within your workload are protected against unintended and unauthorized access and potential vulnerabilities. For example, you'll define trust boundaries (e.g., network boundaries and packet filtering), system security configuration and maintenance (e.g., hardening and patching), operating system authentication and authorizations (e.g., users, keys, and access levels), and other appropriate policy-enforcement points (e.g., web application firewalls and/or API gateways).

In AWS, there are a number of approaches to infrastructure protection. The following sections describe how to use these approaches:

- Protecting network and host-level boundaries
- System security configuration and maintenance

- Enforcing service-level protection

Protecting Network and Host-Level Boundaries

The careful management of your network topology and design forms the foundation of how you provide isolation and boundaries for resources within your environment. Because resources placed within your environment inherit the security properties of the underlying network, it is critical to establish an appropriate network design for the workload that ensures that only the desired network paths and routing are allowed. This can be done by leveraging multiple layers of protection to provide redundancy for the controls and mitigate the impact of a single layer misconfiguration that could allow inappropriate access.

When defining the network topology, consider which components of the system need to be public, for example, a customer-facing load balancer. Additionally, when you design connectivity, consider whether you need connectivity between your data center and AWS over a private network. Apply appropriate configurations to your virtual private cloud (VPC), subnets, routing tables, network access control lists (NACLs), gateways, and security groups to achieve the network routing as well as host-level protection.

A VPC created using Amazon Virtual Private Cloud (Amazon VPC) allows you to define your network topology that spans an AWS Region with a private IP address range you determine. Within a VPC, you can create subnets in an Availability Zone. Each subnet has an associated route table that defines routing rules for managing the paths that traffic within the subnet takes. You can define a publicly routable subnet by having a route that goes to an internet gateway attached to the VPC. The absence of a route to the internet gateway prevents instances from being directly reachable. A subnet can also have a NACL attached to it (stateless firewall). You can configure the NACL to narrow the scope of traffic allowed. For example, you could only allow the port for the database engine for a subnet hosting the databases.

When a host is launched within a VPC, it has its own security group (stateful firewall). This firewall is outside the operating system layer and can be used to define rules for allowed traffic. You can also define relationships between security groups. For example, instances within a database tier security group only accept traffic from instances within the application tier.

When implementing a VPC design, keep in mind some key considerations for the IP address range that you choose for the VPC. Try to use non-overlapping IP addresses with your other VPCs or your data center. When designing NACL rules, consider that it's a stateless firewall and, therefore, both outbound and inbound rules need to be defined to meet your needs.

Key AWS Services

The key AWS service that supports the protection of network- and host-level boundaries is Amazon VPC. This service gives you the ability to create your private virtual network on AWS. The following services and features are also important:

- **Amazon VPC Security Groups** provide a per-host stateful firewall, allowing you to specify traffic rules and define relationships to other security groups.
- **AWS Shield** is a managed distributed denial of service (DDoS) protection service that safeguards web applications running on AWS.
- **AWS WAF** is a web application firewall that helps protect your web applications from common web exploits that could affect application availability, compromise security, or consume excessive resources.
- **AWS Firewall Manager** is a security management service that makes it easier to centrally configure and manage AWS WAF rules across your accounts and applications.
- **AWS Direct Connect** allows you the ability to establish your own direct connectivity from your data center to your VPC.

Resources

Refer to the following resources to learn more about AWS best practices for protecting network- and host-level boundaries.

Video

- [Advanced VPC Design and New Capabilities for Amazon VPC](#)

Documentation

- [Amazon VPC Documentation](#)

- [Getting started with AWS WAF](#)
- [Network Access Control Lists](#)
- [Security Groups for Your VPC](#)
- [Recommended Network ACL Rules for Your VPC](#)
- [AWS Firewall Manager](#)

System Security Configuration and Maintenance

The careful management of the security configurations of the systems running within your environment forms the foundation of how you will maintain robust, secure, scalable systems. The security posture of your systems is a function of the controls that are available and your own controls such as operating system installed threat detection, CVE and vulnerability scanners, anti-malware detection, and any tools to help verify and maintain the integrity of your operating systems. These controls also form another layer in your defense in-depth strategy.

When you define the approach for securing your system, consider the level of access needed for your system and take a least-privilege approach (for example, only open the ports needed for communication, ensure that the operating system is hardened and that unnecessary tools and permissive configurations are disabled). You should apply appropriate configurations to your operating systems, including strong authorization mechanisms.

You should automate deployments and maintenance, and remove operator access to reduce your surface area. Much of this can be achieved using AWS Systems Manager features for EC2 management, and AWS CloudFormation for infrastructure. You should assess your baseline security exposures and perform routine vulnerability assessments when updates or deployments are pushed. CVE-scanning tools such as Amazon Inspector can be used to help with this and centralize security findings for analysis and remediation. You will want to ensure that your operating system and application configurations, such as firewall settings and anti-malware definitions, are correct and up-to-date. You can use AWS Systems Manager to define and maintain consistent operating system configurations. Use AWS Systems Manager to collect and query configuration about your instances and installed software. Leverage automated patching tools such as AWS Systems Manager Patch Manager to help you

deploy operating system and software patches automatically across large groups of instances.

Key AWS Services

The key AWS features that support the protection of your systems is AWS Systems Manager. The following services and features are also important:

- **Amazon Inspector** can be used to identify vulnerabilities or deviations from best practices in your guest operating systems and applications.
- **AWS CloudFormation** can be used to create and manage infrastructure.

Resources

Refer to the following resources to learn more about AWS best practices for system security configuration and maintenance.

Documentation

- [AWS Systems Manager](#)
- [Amazon Inspector](#)
- [Writing your own AWS Systems Manager documents](#)
- [Replacing a Bastion Host with Amazon EC2 Systems Manager](#)

Enforcing Service-Level Protection

The careful management of the security configurations of AWS service endpoints forms the foundation of how you will maintain secure and authorized access to these endpoints. This level of security is required to ensure that users and automated systems have exactly the level of access needed to perform their tasks (least privilege).

You can protect AWS service endpoints by defining policies using IAM. IAM can help you define policies for access to services and operations. However, for some services, you can also define fine-grained controls to specific resources within those services. Additionally, some resources have their own resource-level policies. For example, Amazon S3 has bucket policies to define access levels to objects and/or entire buckets, and AWS KMS has policies to define

administrators and users of the keys in AWS KMS. Use IAM and the resource policies to define a robust protection scheme for your resources.

When defining a service-level protection approach, ensure that you apply a least-privilege methodology and set service-level access policies accordingly.

Key AWS Services

The key AWS service that supports service-level protection is IAM, which lets you define specific policies for many AWS resources. The following services are also important:

- **AWS KMS** allows you to set policies on the individual key.
- **Amazon S3** allows you to set bucket policies for each S3 bucket.
- **Amazon Simple Notification Service (Amazon SNS)** topics allows you to set a policy statement on each topic

Resources

Refer to the following resources to learn more about AWS best practices for enforcing service-level protection.

Documentation

- [IAM Best Practices](#)
- [Using Key Policies in AWS KMS](#)
- [Using Bucket Policies and User Policies](#)
- [SNS Access Control](#)

Data Protection

Before architecting any system, foundational practices that influence security should be in place. For example, data classification provides a way to categorize organizational data based on levels of sensitivity, and encryption protects data by way of rendering it unintelligible to unauthorized access. These methods are important because they support objectives such as preventing financial loss or complying with regulatory obligations.

In AWS, there are a number of different approaches to consider when addressing data protection. The following section describes how to use these approaches:

- Data classification
- Encryption/tokenization
- Protecting data at rest
- Protecting data in transit
- Data backup/replication/recovery

Data Classification

Data classification provides a way to categorize organizational data based on levels of sensitivity. This includes understanding what data types are available, where the data is located, access levels, and protection of the data (for example, through encryption access control). By carefully managing an appropriate data classification system along with each workload's level of protection requirements, you can map the controls and level of access/protection appropriate to the data. For example, public-facing content is available for anyone to access, whereas important content is encrypted and stored in a protected manner that requires authorized access to a key for decrypting the content.

By using resource tags, IAM policies, AWS KMS, and AWS CloudHSM, you can define and implement your policies for data classification. For example, if you have S3 buckets that contain highly critical data or EC2 instances that process confidential data, they can be tagged with a “DataClassification=CRITICAL” tag. You can define levels of access to the AWS KMS encryption keys through key policies to ensure that appropriate services have access to the sensitive content through a secure mechanism.

When you consider a data classification methodology, balance usability versus access. Also consider the multiple levels of access and nuances for implementing a secure, but still usable, approach for each level. Always consider a defense-in-depth approach and reduce human access to data. For example, require users to strongly authenticate to an application. In addition, ensure that users come from a trusted network path and require access to the decryption keys. Use tools

such as dashboards or automated reporting to give users information from the data rather than giving them direct access to the data.

Key AWS Services

The key AWS feature that supports data classification is resource tagging, which provides an ability to apply custom-defined tags for resources. The following services are also important:

- **AWS KMS** allows you to define encryption keys, encrypt data, and protect keys with IAM and access policies.

Resources

Refer to the following resources to learn more about AWS best practices for data classification.

Documentation

- [Tagging Your Amazon EC2 Resources](#)
- [Amazon S3 Object Tagging](#)
- [Using Key Policies in AWS KMS](#)
- [Using Bucket Policies and User Policies](#)

Encryption/Tokenization

Encryption and tokenization are two important but distinct data protection schemes. *Tokenization* is a process that allows you to define a token to represent an otherwise sensitive piece of information (for example, a token to represent a customer's credit card number). A token must be meaningless on its own. *Encryption* is a way of transforming content in a manner that makes it unreadable without a secret key necessary to decrypt the content back into plain text. Both tokenization and encryption can be used to secure and protect information as appropriate.

By carefully defining your tokenization approach, you can provide additional protection for your content, and you can ensure that you meet your compliance requirements. For example, you can narrow the scope of a credit card processing system if you leverage a token instead of a credit card number. You can define your own tokenization scheme by creating a look-up table in an

encrypted Amazon Relational Database Service (Amazon RDS) database instance or an Amazon DynamoDB table and issue tokens to your applications.

By defining an encryption approach, you can provide protection for your content against unauthorized users and against unnecessary exposure to authorized users. AWS KMS helps you manage encryption keys and integrates with many AWS services. This service provides durable, secure, and redundant storage for your master keys. You can define your key aliases as well as key-level policies. The policies help you define key administrators as well as key users. For example, a secret management system can be the only system that has access to the master key that encrypts the secrets for storage.

Additionally, AWS CloudHSM is a cloud-based hardware security module (HSM) that enables you to easily generate and use your own encryption keys on the AWS Cloud. It helps you meet corporate, contractual, and regulatory compliance requirements for data security by using FIPS 140-2 Level 3 validated HSMs.

When defining an encryption/tokenization approach, consider the data classification model you've defined and the levels of access needed for each content. Consider the compliance requirements and needs around the content and how to strictly enable that approach. Also carefully consider the differences and different use cases for tokenization versus encryption. Consider the key policies and access levels that would be provided for the user.

Key AWS Services

The key AWS service that supports encryption is AWS KMS, which provides an easy-to-use, secure, and redundant key management service. The following services are also important:

- **AWS CloudHSM** provides a hardware security module for managing your keys
- **Amazon DynamoDB** provides a way to implement a fast NoSQL database. This can be used to store encrypted content for your tokens.

Resources

Refer to the following resources to learn more about AWS best practices for encryption/tokenization.

Video

- [Best Practices for Implementing AWS Key Management Service](#)
- [A Deep Dive into AWS Encryption Services](#)

Documentation

- [Protecting Data Using Encryption](#)
- [Amazon EBS Encryption](#)
- [AWS Key Management Service](#)
- [AWS CloudHSM](#)
- [AWS KMS Cryptographic Details Whitepaper](#)

Protecting Data at Rest

Data at rest represents any data that you persist for any duration. This includes block storage, object storage, databases, archives, and any other storage medium on which data is persisted. Protecting your data at rest reduces the risk of unauthorized access, when encryption and appropriate access controls are implemented.

Multiple AWS services provide built-in integration with AWS KMS to allow easy encryption of your data. Amazon S3 allows you to encrypt content by selecting a KMS key on object upload. Amazon Elastic Block Store (Amazon EBS) allows you to choose a KMS key to encrypt a block storage volume or Amazon Machine Image (AMI) copy operation. Amazon RDS allows you to choose an encryption key for encrypting DB instance storage at rest (including backup snapshots).

You also have the option of implementing your own encryption-at-rest approach. For example, you can encrypt content before storing in an AWS service. Amazon S3 provides you the facility to upload an already encrypted object. It also provides the ability for you to upload an object along with an encryption key that's used in-memory to encrypt an object. To retrieve the object, you must supply the same key.

When implementing an encryption-at-rest protection method, consider the data classification model for your organization to ensure that the content's protection reflects your business, legal, compliance, and regulatory requirements.

Finally and most importantly, ensure that you have implemented a least-privileged approach to control access to the keys, storage mediums, and any compute resources that have access to the content.

Key AWS Services

The key AWS service that protects data at rest is AWS KMS, which is a managed service that makes it easy for you to create and control encryption keys. The following services are also important:

- **Amazon S3** is an object storage service that integrates with AWS KMS, and allows you to supply your own keys.
- **Amazon EBS** is a block storage service integrated with AWS KMS. You can also perform block-level encryption with your operating system tools or third-party solutions.
- **Amazon Glacier** is a secure, durable, and extremely low-cost storage service for data archiving and long-term backup that encrypts data at rest.

Resources

Refer to the following resources to learn more about AWS best practices for protecting data at rest.

Video

- [Encryption and Key Management in AWS](#)

Documentation

- [Protecting Amazon S3 Data Using Encryption](#)
- [Amazon EBS Encryption](#)
- [How Amazon Simple Storage Service \(Amazon S3\) Uses AWS KMS](#)
- [How Amazon Elastic Block Store \(Amazon EBS\) Uses AWS KMS](#)
- [Encrypting Amazon RDS Resources](#)

Protecting Data in Transit

Data in transit is any data that gets transmitted from one system to another. This includes communication between resources within your environment as well as communication between other services and your end users. By providing the appropriate level of protection for your data in transit, you protect the confidentiality and integrity of your application's data. When protecting your data in transit, selecting secure protocols that implement the latest in cryptography standards such as Transport Layer Security (TLS) is a common best practice.

AWS services provide HTTPS endpoints using TLS for communication, thus providing encryption in transit when communicating with the AWS APIs. You have full control over your computing resources to implement encryption in transit across your services. Additionally, the AWS Certificate Manager (ACM) service provides you the ability to manage and deploy public and private certificates for your workloads.

Additionally, you can leverage VPN connectivity into your VPC or across your VPCs to facilitate encryption of traffic.

When planning for an encryption-in-transit approach, consider your use cases and the balance between encryption and ease of use. Consider the use of VPN connectivity into AWS and look into HTTPS for application-to-application communication in a secure manner. Using Elastic Load Balancing (ELB), Amazon CloudFront, and ACM make it easy to generate, deploy, and manage certificates used for TLS encryption in web-based workloads.

Key AWS Services

The key AWS service that protects data in transit is **ACM**, which helps you generate certificates used for establishing encrypted transport between systems. The following services and features are also important:

- **Elastic Load Balancing** supports secure protocols including HTTPS with ACM integration.
- **Amazon CloudFront** supports encrypted endpoints for your content distributions.

Resources

Refer to the following resources to learn more about AWS best practices for protecting data in transit.

Documentation

- [AWS Certificate Manager](#)
- [HTTPS Listeners for Your Application Load Balancer](#)
- [VPN Connections](#)

Data Backup/Replication/Recovery

By defining your data backup, replication, and recovery approach, you help protect against deletion or destruction of data. A sound approach for data backup and replication helps protect you in case of a disaster. Properly secured and protected primary and secondary data sources can ensure continued business operations.

AWS provides you with multiple features and capabilities for data backup and replication. Amazon S3 is designed for 99.999999999% durability for objects stored in the service, and it allows you to create copies of the content that can be replicated to other locations and AWS accounts for additional protection. Amazon RDS performs snapshots of your database instances and allows you to replicate those instances to other locations. Snapshots can be taken of EBS volumes and copied across AWS Regions if required. Additionally, you can automate tasks and scheduled jobs to perform backups of resources. Consider using a Lambda function to issue backups at an interval using CloudWatch Events.

Amazon Glacier is a secure, durable, and extremely low-cost cloud storage service for data archiving and long-term backup. Use this service for cost-effective storage of backup copies. These copies can be retrieved when needed for recovery, whether it's for testing, regulatory reasons, or disaster scenarios.

When defining your backup/replication/recovery approach, ensure that you review the scenarios under which you want to be protected as well as the nuances of each. For example, Amazon RDS will replicate any accidental changes, so if you have backups from multiple points in time you may be

protected from accidental errors in addition to malicious actions. Ensure that you have a process defined for the recovery of your content. You should plan and run a game day scenario to test and ensure that your approach is effective in the event of a disaster. Moreover, consider storing backups in a different AWS account with a different set of credentials to protect against human error or a compromise of the primary account.

Key AWS Services

The key AWS service that supports data backup, replication, and recovery is **Amazon S3**. The following services and features are also important for backup and replication of data:

- **Amazon S3 Cross-Region Replication** is an Amazon S3 bucket-level feature that enables automatic, asynchronous copying of objects across buckets in different AWS Regions.
- **Amazon S3 lifecycle policies and versioning** allow you to implement a backup strategy and meet retention requirements.
- **Amazon EBS snapshot operations** let you back up your volumes attached to EC2 instances.

Resources

Refer to the following resources to learn more about AWS best practices for data backup, replication, and recovery.

Documentation

- [Amazon S3 Cross-Region Replication](#)
- [Amazon S3 Object Lifecycle Management](#)
- [Amazon S3 Object Versioning](#)
- [Amazon EBS Snapshots](#)
- [Getting Started with Amazon Glacier](#)
- [Using AWS Lambda with Scheduled Events](#)

Incident Response

Even with extremely mature preventive and detective controls, your organization should still put processes in place to respond to and mitigate the potential impact of security incidents. The architecture of your workload strongly affects the ability of your teams to operate effectively during an incident, to isolate or contain systems, and to restore operations to a known good state. Putting in place the tools and access ahead of a security incident, then routinely practicing incident response through game days, will help you ensure that your architecture can accommodate timely investigation and recovery.

In AWS, there are a number of different approaches to consider when addressing incident response. The Clean Room section describes how to use these approaches.

Clean Room

In every incident, maintaining situational awareness is one of the most important principles. By using tags to properly describe your AWS resources, incident responders can quickly determine the potential impact of an incident. For example, tagging instances and other resources with an owner or work queue in a ticketing system allows the team to engage the right people more quickly. By tagging systems with a data classification or a criticality attribute, the impact of an incident can be estimated more accurately.

During an incident, the right people require access to isolate and contain the incident, and then perform forensic investigation to identify the root cause quickly. In some cases, the incident response team is actively involved in remediation and recovery as well. Determining how to get access for the right people during an incident delays the time it takes to respond, and can introduce other security weaknesses if access is shared or not properly provisioned while under pressure. Determine the access your team members need ahead of time, and then regularly verify the access is functional - or easily triggered—when needed.

In AWS you can use the power of the APIs to automate many of the routine tasks that need to be performed during an incident and subsequent investigations. For example, you can isolate an instance by changing the

security groups associated with an instance or removing it from a load balancer. Architecting your workload using Auto Scaling potentially allows the instance under investigation to be removed from production without affecting the availability of your applications.

Forensics often requires capturing the disk image or "as-is" configuration of an operating system, you can use EBS snapshots and the Amazon EC2 APIs to capture the data and state of systems under investigation. Storing snapshots and related incident artifacts in Amazon S3 ensures that the data will be available and retained as appropriate.

During an incident, before the root cause has been identified and the incident has been contained, it can be difficult to conduct investigations in an untrusted environment. Unique to AWS, security practitioners can use CloudFormation to quickly create a new, trusted environment in which to conduct deeper investigation. The CloudFormation template can pre-configure instances in an isolated environment that contains all the necessary tools forensic teams need to determine the cause of the incident. This cuts down on the time it takes to gather necessary tools, isolate systems under examination, and ensures that the team is operating in a clean room.

Key AWS Services

Several key AWS services and features are critical to a mature incident response process:

- **IAM** should be used to grant appropriate authorization to incident response teams in advance.
- **AWS CloudFormation** to automate the creation of trusted environments for conducting deeper investigations.
- **AWS CloudTrail** provides a history of AWS API calls that can assist in response, and trigger automated detection and response systems.
- **Amazon CloudWatch Events** to trigger different automated actions from changes in AWS resources including CloudTrail.
- **AWS Step Functions** to coordinate a sequence of steps to automate an incident response process.

Resources

Refer to the following resources to learn more about AWS best practices for incident response.

Videos

- [Automating Incident Response and Forensics](#)
- [Incident Response in the Cloud](#)

Documentation

- [Incident Response on AWS](#)
- [AWS Step Functions](#)
- [AWS CloudTrail](#)
- [Amazon CloudWatch Events](#)

Conclusion

Security is an ongoing effort. When incidents occur they should be treated as opportunities to improve the security of the architecture. Having strong authentication and authorization controls, automating responses to security events, protecting infrastructure at multiple levels, and managing well-classified data with encryption provides defense-in-depth that every business should expect. This effort is easier thanks to the programmatic functions and AWS features and services discussed in this paper.

AWS strives to help you build and operate architectures that protect information, systems, and assets while delivering business value. To make your architectures truly secure, you should use the tools and techniques discussed in this paper.

Contributors

The following individuals and organizations contributed to this document:

- Philip Fitzsimons, Sr Manager Well-Architected, Amazon Web Services
- Bill Shinn, Principal, Office of the CISO, Amazon

- Sam Elmalak, Solutions Architect, Amazon Web Services
- Ben Potter, Security Lead Well-Architected, Amazon Web Services

Further Reading

For additional help, please consult the following sources:

- [AWS Well-Architected Framework Whitepaper](#)

Document Revisions

Date	Description
July 2018	Updates to reflect new AWS services and features, and updated references.
May 2017	Updated System Security Configuration and Maintenance section to reflect new AWS services and features.
November 2016	First publication